

Assignment Week11

May 24, 2021

```
[374]: import time
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

import pandas as pd

from sklearn.datasets import fetch_openml
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.utils import check_random_state
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix, plot_confusion_matrix, roc_curve, auc, precision_score, \
    recall_score
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

```
[21]: df =pd.read_csv("Data_for_UCI_named.csv")
df
```

```
[21]:
```


	tau1	tau2	tau3	tau4	p1	p2	p3	\
0	2.959060	3.079885	8.381025	9.780754	3.763085	-0.782604	-1.257395	
1	9.304097	4.902524	3.047541	1.369357	5.067812	-1.940058	-1.872742	
2	8.971707	8.848428	3.046479	1.214518	3.405158	-1.207456	-1.277210	
3	0.716415	7.669600	4.486641	2.340563	3.963791	-1.027473	-1.938944	
4	3.134112	7.608772	4.943759	9.857573	3.525811	-1.125531	-1.845975	
...	
9995	2.930406	9.487627	2.376524	6.187797	3.343416	-0.658054	-1.449106	
9996	3.392299	1.274827	2.954947	6.894759	4.349512	-1.663661	-0.952437	
9997	2.364034	2.842030	8.776391	1.008906	4.299976	-1.380719	-0.943884	
9998	9.631511	3.994398	2.757071	7.821347	2.514755	-0.966330	-0.649915	
9999	6.530527	6.781790	4.349695	8.673138	3.492807	-1.390285	-1.532193	
	p4	g1	g2	g3	g4	stab	stabf	

0	-1.723086	0.650456	0.859578	0.887445	0.958034	0.055347	unstable
1	-1.255012	0.413441	0.862414	0.562139	0.781760	-0.005957	stable
2	-0.920492	0.163041	0.766689	0.839444	0.109853	0.003471	unstable
3	-0.997374	0.446209	0.976744	0.929381	0.362718	0.028871	unstable
4	-0.554305	0.797110	0.455450	0.656947	0.820923	0.049860	unstable
...
9995	-1.236256	0.601709	0.779642	0.813512	0.608385	0.023892	unstable
9996	-1.733414	0.502079	0.567242	0.285880	0.366120	-0.025803	stable
9997	-1.975373	0.487838	0.986505	0.149286	0.145984	-0.031810	stable
9998	-0.898510	0.365246	0.587558	0.889118	0.818391	0.037789	unstable
9999	-0.570329	0.073056	0.505441	0.378761	0.942631	0.045263	unstable

[10000 rows x 14 columns]


```
[44]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.distplot(df.tau1, hist = False)
```

[44]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb6d2a6af40>



```
[322]: figure = plt.figure(figsize = (10,10))
column_list = list(df.columns[:-3])
for i, v in enumerate(column_list):
    figure.add_subplot(3, 4, i+1)
```

```
sns.distplot(df[v], hist = False)
```



```
[51]: df.stabf.value_counts()
```

```
[51]: unstable    6380
      stable      3620
      Name: stabf, dtype: int64
```

```
[119]: df['Target'] = [1 if x== 'stable' else 0 for x in df.stabf]
```

```
[150]: df.Target.value_counts()
```

```
[150]: 0    6380
      1    3620
      Name: Target, dtype: int64
```

```
[190]: X = df.iloc[:, :-3]
      y= df['Target']
```

```
[191]: X_train, X_test, y_train, y_test = train_test_split(X,y, train_size= 0.8)
```

```
[192]: from sklearn.preprocessing import StandardScaler
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[289]: lr = LogisticRegression( solver = "saga")
```

```
[290]: parameters = [{'C':[0.001, 0.01, 0.1,1], 'tol':[0.001, 0.01,0.05,0.1],
      ↪'penalty':['l1','l2']}]
```

```
[291]: grid1 =GridSearchCV(lr, parameters, scoring = 'roc_auc', cv=4)
```

```
[292]: grid1.fit(X, y)
```

```
[292]: GridSearchCV(cv=4, estimator=LogisticRegression(solver='saga'),
      param_grid=[{'C': [0.001, 0.01, 0.1, 1], 'penalty': ['l1', 'l2'],
      'tol': [0.001, 0.01, 0.05, 0.1]}],
      scoring='roc_auc')
```

```
[293]: print("best parameter is: %s"%grid1.best_params_)
```

```
best parameter is: {'C': 1, 'penalty': 'l2', 'tol': 0.001}
```

```
[294]: grid1.best_score_ # highest AUC value
```

```
[294]: 0.8906008070801366
```

```
[295]: grid1.best_index_
```

```
[295]: 28
```

```
[296]: grid1.cv_results_["mean_test_score"]
```

```
[296]: array([0.78912382, 0.78913663, 0.78908537, 0.78895772, 0.81218137,
      0.81217461, 0.81172448, 0.81056184, 0.87472661, 0.87307799,
      0.86232737, 0.84927657, 0.87070178, 0.86938413, 0.8472102 ,
      0.83976047, 0.89054781, 0.88992951, 0.88083029, 0.8659594 ,
      0.89003776, 0.88908 , 0.87506157, 0.86329708, 0.89059561,
```

```
0.88996467, 0.87920331, 0.87036925, 0.89060081, 0.88993852,  
0.87971787, 0.86771662])
```

```
[297]: grid1.cv_results_["std_test_score"][28]
```

```
[297]: 0.011710624735358218
```

```
[298]: print("average of auc across 4 folds for the best parameter: %.2f, with a std_  
      ↪dev of : %.2f"  
      %(grid1.best_score_, grid1.cv_results_["std_test_score"][28]))
```

```
average of auc across 4 folds for the best parameter: 0.89, with a std dev of :  
0.01
```

```
[299]: lr = LogisticRegression(C = 1, penalty = 'l2', solver = 'saga', tol = 0.001) #_  
      ↪Fitting a logistic Regression Model
```

```
[300]: lr.fit(X_train, y_train)  
      y_predict = lr.predict(X_test)
```

1 Model 1 Logistic Regression Evaluation:

```
[301]: precision_score(y_test, y_predict, average = 'weighted') # using weighted_  
      ↪because we have an imbalanced target data
```


```
[301]: 0.8031790667300063
```

```
[302]: recall_score(y_test, y_predict, average = "weighted")
```

```
[302]: 0.8055
```

```
[303]: plot_confusion_matrix(lr, X_test, y_test)
```

```
[303]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
      0x7fb6d29c04f0>
```



```
[304]: from sklearn.model_selection import cross_val_score
score = cross_val_score(lr, X,y, scoring='roc_auc', cv=4)
```

```
[307]: score.mean() # with cross validation, auc score is 0.89
```

```
[307]: 0.8906020194322728
```


```
[308]: roc_auc_score(y_test, lr.predict_proba(X_test)[:,:1]) #with train, test split,
↪method, auc score is 0.88
```

```
[308]: 0.8827246251752777
```

```
[309]: fp_lr, tp_lr, _ = roc_curve(y_test, lr.predict_proba(X_test)[:,:1], pos_label = 1)
```

```
[310]: plt.plot(fp_lr, tp_lr)
```

```
[310]: [<matplotlib.lines.Line2D at 0x7fb6d242dbb0>]
```



2 Model 2 Random Forest Classifier

```
[275]: rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
y_predict_rf = rfc.predict(X_test)
```

```
[276]: roc_auc_score(y_test, rfc.predict_proba(X_test)[:,-1])
```

```
[276]: 0.9790211411929673
```

```
[278]: precision_score(y_test, y_predict_rf, average='weighted') # using weighted
      ↪ because we have an imbalanced target data
```


```
[278]: 0.92028159165258
```

```
[279]: recall_score(y_test, y_predict_rf, average = "weighted")
```

```
[279]: 0.9205
```

```
[281]: plot_confusion_matrix(rfc, X_test, y_test)
```

```
[281]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb6d263cd90>
```



```
[280]: fp_rf, tp_rf, _ = roc_curve(y_test, rfc.predict_proba(X_test)[: ,1])
```

3 Model 3 KNN Classifier

```
[334]: from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
```

```
[329]: knn = KNeighborsClassifier()
```


```
[335]: knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)
yhat = knn.predict(X_test)
metrics.accuracy_score(y_test, yhat)
```

```
[335]: 0.849
```

```
[336]: acc_score = []
for k in range(1,100):
    knn = KNeighborsClassifier(n_neighbors = 3)
    knn.fit(X_train, y_train)
    yhat = knn.predict(X_test)
    acc_score.append(metrics.accuracy_score(y_test, yhat))
```



```
[347]: plt.plot(acc_score); # n value does not matter!
```



```
[345]: roc_auc_score(y_test, knn.predict_proba(X_test)[: ,1])
```

```
[345]: 0.9017614065365117
```

```
[348]: precision_score(y_test, yhat, average = 'weighted') # using weighted because we have an imbalanced target data
```


```
[348]: 0.8479407372799063
```

```
[351]: recall_score(y_test, yhat, average = 'weighted') # using weighted because we have an imbalanced target data
```

```
[351]: 0.849
```

```
[353]: plot_confusion_matrix(knn, X_test, y_test)
```

```
[353]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fb6d23df070>
```



```
[354]: fp_knn, tp_knn, _ = roc_curve(y_test, knn.predict_proba(X_test)[: ,1])
```

4 Model 4 Decision Tree Classifier

```
[358]: DT = DecisionTreeClassifier()
```

```
[362]: parameters = [{"criterion":["entropy","min_weight_fraction_leaf","poisson"],
    ↪ "max_depth":list(np.arange(10))}]
grid4 = GridSearchCV(DT,parameters , scoring = 'roc_auc', cv=4)
```

```
[ ]: grid4.fit(X,y)
```

```
[365]: grid4.best_params_
```

```
[365]: {'criterion': 'entropy', 'max_depth': 7}
```

```
[366]: DT = DecisionTreeClassifier(criterion ="entropy", max_depth =7)
```

```
[367]: DT.fit(X_train, y_train)
```

```
[367]: DecisionTreeClassifier(criterion='entropy', max_depth=7)
```

```
[368]: y_predict_dt = DT.predict(X_test)
```

```
[369]: roc_auc_score(y_test, DT.predict_proba(X_test)[: ,1])
```

```
[369]: 0.8937477079063747
```

```
[370]: precision_score(y_test, y_predict_dt, average = 'weighted') # using weighted  
      ↪ because we have an imbalanced target data
```


```
[370]: 0.8235872335240983
```

```
[371]: recall_score(y_test, y_predict_dt, average = 'weighted') # using weighted  
      ↪ because we have an imbalanced target data
```

```
[371]: 0.8255
```

```
[372]: plot_confusion_matrix(DT, X_test, y_test)
```

```
[372]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
      0x7fb6d2291a00>
```



```
[373]: fp_DT, tp_DT, _ = roc_curve(y_test, DT.predict_proba(X_test)[: ,1])
```

5 Model 5 Gradient Boosting Classifier

```
[375]: gb = GradientBoostingClassifier(n_estimators=100, max_depth=1)
      params = [{'n_estimators':[50, 75,100,125,150], "max_depth":[1,2,3,4,5],
      ↪ "max_features": ['auto', 'sqrt', 'log2']}]

[376]: grid4 =GridSearchCV(gb,params , scoring = 'roc_auc', cv=4)

[377]: grid4.fit(X,y)

[377]: GridSearchCV(cv=4, estimator=GradientBoostingClassifier(max_depth=1),
      param_grid=[{'max_depth': [1, 2, 3, 4, 5],
      'max_features': ['auto', 'sqrt', 'log2'],
      'n_estimators': [50, 75, 100, 125, 150]}],
      scoring='roc_auc')

[378]: grid4.best_params_

[378]: {'max_depth': 5, 'max_features': 'auto', 'n_estimators': 150}

[383]: grid4.best_score_ # The highest ROC_AUC score yet!

[383]: 0.9866899322814735

[379]: gb = GradientBoostingClassifier(n_estimators=150, max_depth=5, max_features_
      ↪="auto")

[380]: gb.fit(X_train, y_train)

[380]: GradientBoostingClassifier(max_depth=5, max_features='auto', n_estimators=150)

[381]: y_predict_gb = gb.predict(X_test)

[384]: precision_score(y_test, y_predict_gb , average ='weighted') # using weighted_
      ↪because we have an imbalanced target data


[384]: 0.9383973426405792

[385]: recall_score(y_test, y_predict_gb , average ='weighted') # using weighted_
      ↪because we have an imbalanced target data

[385]: 0.9385


[386]: plot_confusion_matrix(gb, X_test, y_test)

[386]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
      0x7fb6d2351ac0>
```



```
[388]: fp_gb, tp_gb, _ = roc_curve(y_test, gb.predict_proba(X_test)[: ,1])
```

```
[393]: plt.plot(fp_lr, tp_lr, label = "Logistic Regression")
plt.plot(fp_rf, tp_rf, label = "Random Forest")
plt.plot(fp_knn, tp_knn, label = "K Nearest Neighbor")
plt.plot(fp_DT, tp_DT, label = "Decision Trees")
plt.plot(fp_gb, tp_gb, label = "Gradient Boosting")
plt.legend()
plt.title("ROC Curve of FIVE MODELS");
```



6 Gradient Boosting outperforms all others!

[]: